

- Target

The target will be identifying negative reviews. It will be created based on the `Rating` column, and will be named `Sentiment`. Based on a given set of reviews, I will predict if the review's sentiment was negative. To be more precise, it will be identified as `detractors`. This term is defined by the metric measuring online reputation: Net Promoter Score.

More on NPS below.

4. Data Preparation

4.1- Data Cleaning

In the first part of data preparation, the typical data cleaning tasks are addressed before splitting the set between train and test data. The steps include:

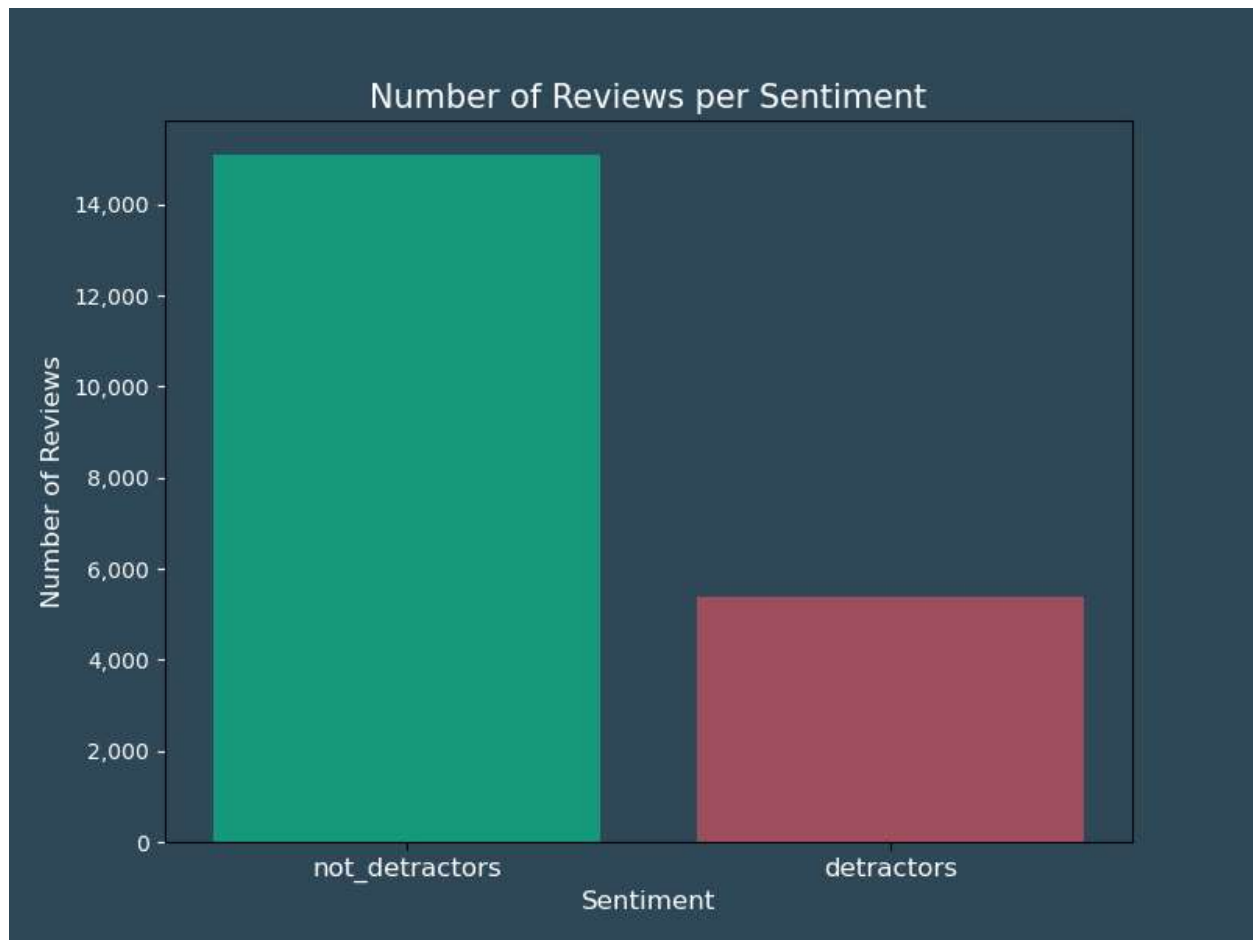
- Missing data
Missing values were inspected but none were found, no modification was necessary.
- Duplicates
Duplicates were researched but none were found: no modification was necessary.
- Binary classification
Ratings range from 1 to 5, 1 being the lowest and 5 the highest.
A new column `Sentiment` was created to group customers' reviews as follows:
 - Reviews' rating ranging from 1 to 3 included: `detractors`
 - All other reviews: `not_detractors`

The reason for such a classification follows the definition of NPS score.

NPS Score



The Net Promoter Score ([NPS](#)) is a metric used to measure customer satisfaction and loyalty with a hotel, (or more generally a product, service, or company). It is based on the question: "How likely is it that you would recommend our hotel? The responses can be measured on a scale from 0 to 10 or in our case: 1 to 5.



- Train-Test split

The dataset is being divided into two separate subsets: a training set, and a testing (or validation) set. The validation set will allow to assess the performance of the model. The dataset is split before any further transformation is done to prevent data leakage.

Two parameters were assigned:

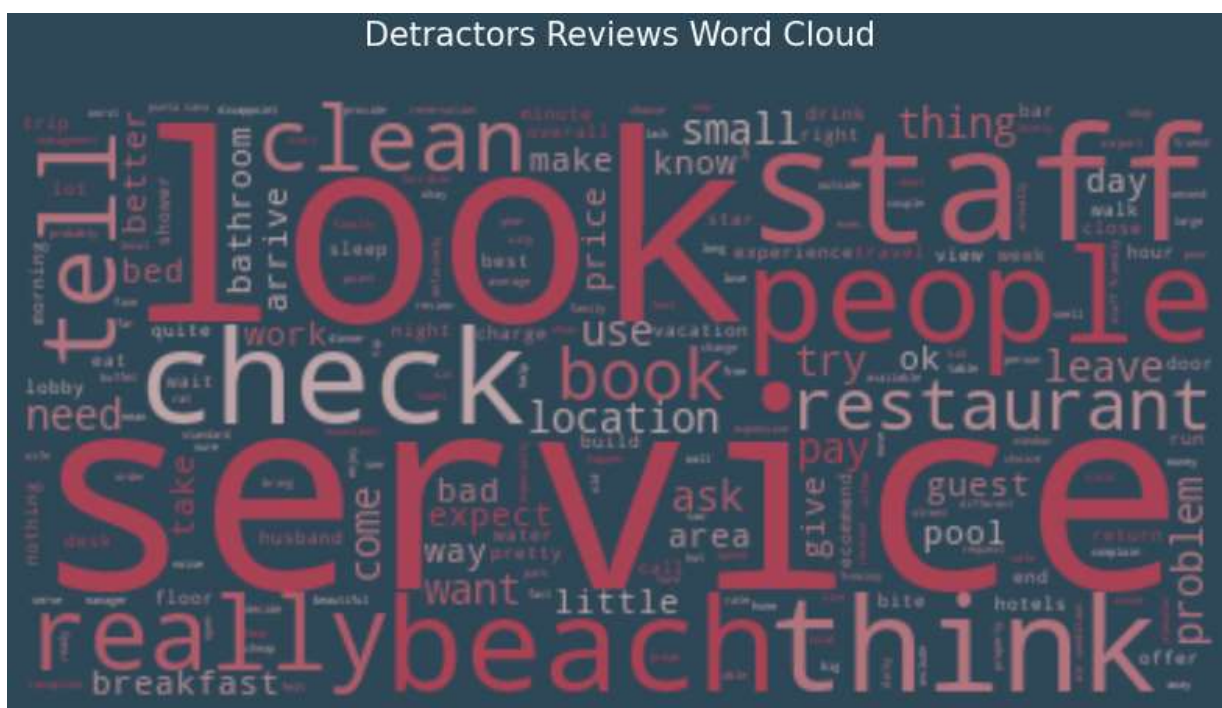
- `random_state=42` for reproducibility
- `stratify=y` to address class imbalance issues
- Encoding target In order to make predictions useable in calculations, I encoded the target:
 - 0 replaced `not_detractors`
 - 1 replaced `detractors`

4. 2- Data Preprocessing & Exploratory Analysis

In order to preprocess the reviews, the following transformations were performed:

- **Standardizing case**
This step was verified as it is important to ensure text is uniform and consistent. Nevertheless, no extra step was taken as text was already saved as lower case.
- **Character Encoding**
This was done to ensure consistent representation of text by transforming non-text characters into a normalized format, which will facilitate proper text processing, analysis, and model training.

- **Tokenizing**
Tokens of words were created. This was done with the `RegexTokenizer` package from `nltk.tokenize`. This step will facilitate the conversion of word into a suitable form for analysis and modeling.
- **Stopwords**
To focus on the data's theme, English stopwords were removed. Manual additions were made in this text's context, such as adding industry specific word (i.e. "hotel", "stay")
- **Lemmatize**
The `WordNetLemmatizer` package from `nltk.stem.wordnet` was used to reduce words to their base form, allowing a more accurate analysis
- **Frequency Distribution**
The `FreqDist` package was used to review in a dictionary-like output, the words and their frequencies
- **WordCloud**
The words' frequencies were represented visually thanks to the `WordCloud` package



- **Bigrams**
Bigrams were drawn to have a better understanding of the themes thanks to the `collocations` package and its `BigramAssocMeasures`
- **Mutual Information Scores**
Bigrams that occur more than 5 times were examined through `mutual information scores`

The preprocessing tasks were summarized as a function which was called both on the train and test data.

5. Modeling

I now have an initial idea for recommendations on focus areas to improve guest satisfaction.

My objective is now to:

- Develop a tool to identify unsatisfied customer reviews in real time

Incorrectly identifying a guest as satisfied when they weren't would lead in a detractor posting a review instead of catching them while they are in-house, and having a chance of improving their satisfaction.

As a consequence, `recall` is the main evaluation metric for this model.

As the dataset is a text, it requires a transformation before it can be used for modeling. Like other types of dataset would one-hot encoded, here, the reviews were vectorized, using the common method in natural language processing:

`TfidfVectorizer`.

It converts a collection of text documents to a matrix of tf-idf features.

- Term-Frequency
Measures how often a term (word) appears in a document
- Inverse Document Frequency (IDF)
Measures the importance of a term in the entire collection of documents.

7 main classification models were explored:

1. Multinomial Naive Bayes
2. K-Nearest Neighbor
3. Decision Tree
4. Random Forest
5. Gradient Boosting
6. AdaBoost
7. XGBoost


Undersampling, stopwords, lemmatize, and hyperparameter tuning were parameters used to tune models. In addition, `X_train` was split with `stratify` in order to keep the distribution ratios for each sentiment.

All models went through 4 steps: 1) Fitting and training on train data 2) Evaluation Metrics 3) Classification Report 4) Confusion Matrix 5) ROC-AUCs

The 7 models' results and their evaluation metrics are summarized below. The highest records are highlighted.

	models	accuracy	f1	recall	recall_detractors	cross-val_accuracies	ROC-AUCs
0	Baseline	0.754636	0.665579	0.754636	0.06963	0.750716	0.911003
1	Resampled	0.884248	0.878426	0.884248	0.65481	0.883589	0.927415
2	Sampling 1	0.831154	0.837767	0.831154	0.84296	0.844221	0.919755
3	No Stopwords	0.831934	0.838196	0.831934	0.83333	0.846628	0.918632
4	Industry Stopwords	0.833106	0.839580	0.833106	0.84444	0.845977	0.920165
5	Full Stopwords	0.834472	0.840474	0.834472	0.83259	0.848060	0.919093
6	Lemmatized	0.827054	0.834410	0.827054	0.85630	0.838169	0.919214
7	Preprocessed	0.827054	0.834132	0.827054	0.84593	0.840381	0.917364
8	NaiveBayes GS	0.827445	0.834899	0.827445	0.86148	0.836998	0.914550
9	kNN	0.745071	0.756900	0.745071	0.72815	0.756442	0.803643
10	DecisionTree	0.725942	0.738918	0.725942	0.69926	0.726834	0.717374
11	RandomForest	0.833301	0.838787	0.833301	0.81259	0.850272	0.913850
12	GradientBoosting	0.824907	0.831126	0.824907	0.81037	0.837323	0.902521
13	AdaBoost	0.814952	0.822096	0.814952	0.80889	0.827563	0.890777
14	XGBoost	0.846769	0.851977	0.846769	0.84370	0.857366	0.922079
15	Tuned XGBoost	0.851845	0.856469	0.851845	0.83926	0.856715	0.924855

The most relevant ones can be summarized through the below bar chart.



6. Evaluation

6. a) Final Model and Classification Metrics

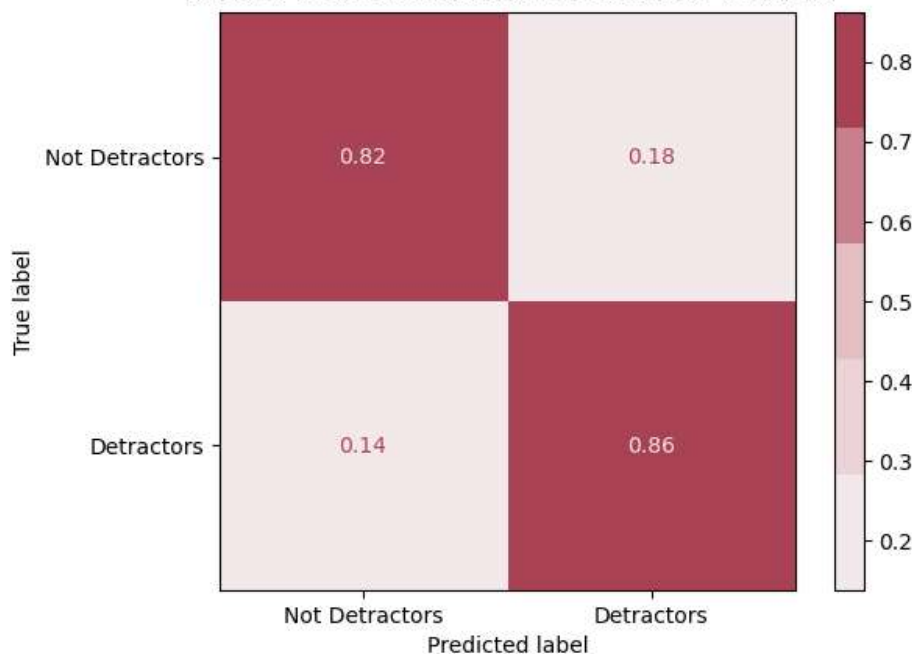
The model that predicts the most accurately the detractors reviews is the **Naive Bayes Tuned with Grid Search**.

Evaluation Metrics

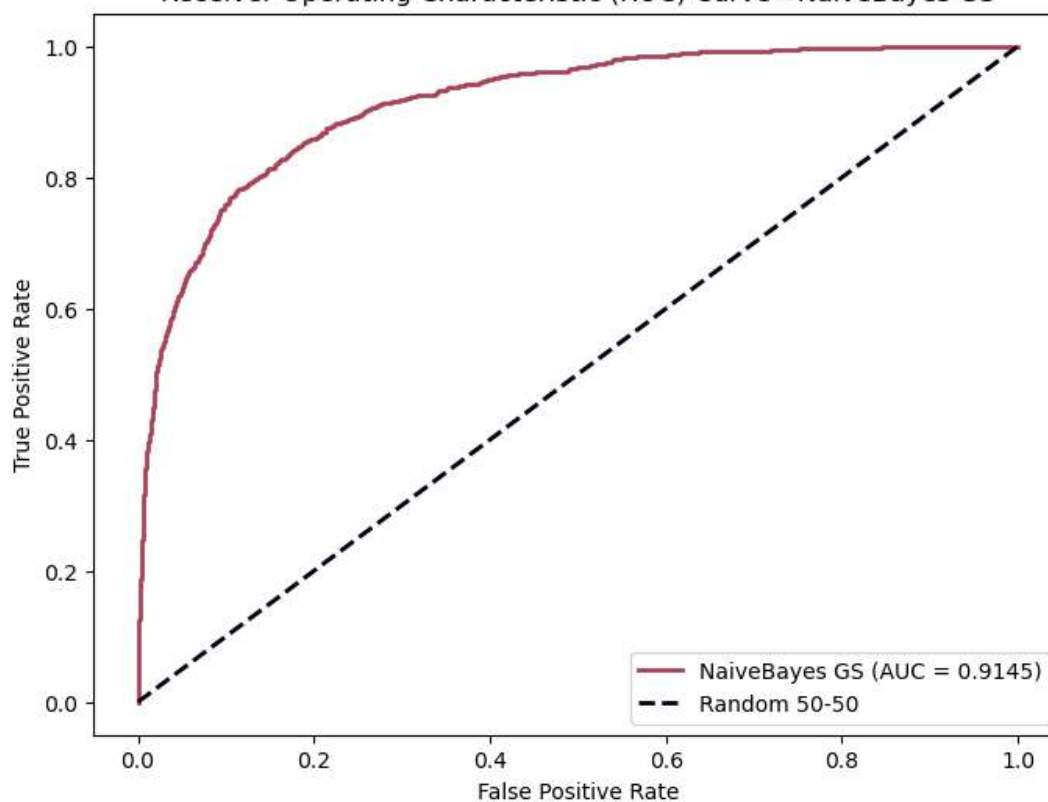
- Evaluation Metrics on Train Data
 - Accuracy: 0.8274
 - Recall Detractors: 0.8615
 - Train data: Mean Cross-Validated Accuracy: 0.8370
- Test data: Mean Cross-Validated Accuracy: 0.8454

6. b) Model Performance

Model Performance: Confusion Matrix



Receiver Operating Characteristic (ROC) Curve - NaiveBayes GS



The classification report, confusion matrix and ROC-AUC summarize the evaluation of the model's performance on predicting detractors from hotel TripAdvisor reviews.

The model performs better on predicting *detractors* sentiment. More precisely, it has a true positive rate higher for this class than for the *not_detractors*.

This score was the focus for the evaluation of the models, as the cost of false negative is higher than the cost of false positive. If a review from a detractor was incorrectly identified as "not detractor", the hotel would miss an opportunity to transform a guest's negative experience.

Looking at the details by metric:

Accuracy

The mode correctly identified close to 83% of all reviews.

Recall

86.15% of actual detractors were correctly predicted.

F1 Score

The balance between precision and recall. It is lower (72.5%) for the detractors class, which indicates precision is lower for this class.

AUC


The Area Under the Receiver Operating Characteristic Curve (AUC) value of 0.9145 indicates a high effectiveness in distinguishing between the two classes: `not_detractors` and `detractors`. The ROC curve, which plots the true positive rate against the false positive rate, covers a significant area under the curve.

7. Findings & Recommendations

7. Recommendations

The advantage of our new service is the deployment can be immediate. Based on historical guest complaints, below are the areas we recommend to focus on to immediately start improving guest satisfactions.

1. Focus on resort hotels
 - These hotels have the most unsatisfied guests
2. Develop a maintenance program with engineering teams
 - Appearance & dysfunctionality cause frustrations
3. Train staff to enhance friendliness
 - Also ensure more languages are spoken
4. Respond to reviews
 - Detractors advise to read reviews



Releases

No releases published
[Create a new release](#)

Packages

No packages published
[Publish your first package](#)

Languages

● Jupyter Notebook 100.0%

